

Network Working Group
Internet-Draft
Intended status: Informational
Expires: August 18, 2014

Z. Sarker
Ericsson AB
V. Singh
Aalto University
X. Zhu
M. A. Ramalho
Cisco Systems
February 14, 2014

Test Cases for Evaluating RMCAT Proposals
draft-sarker-rmcat-eval-test-00

Abstract

The Real-time Transport Protocol (RTP) is used to transmit media in multimedia telephony applications, these applications are typically required to implement congestion control. The RMCAT working group is currently working on candidate algorithms for such interactive real-time multimedia applications. This document describes the test cases to be used in the evaluation of the performance of those candidate algorithms.

Status of This Memo

This Internet-Draft is submitted in full conformance with the provisions of BCP 78 and BCP 79.

Internet-Drafts are working documents of the Internet Engineering Task Force (IETF). Note that other groups may also distribute working documents as Internet-Drafts. The list of current Internet-Drafts is at <http://datatracker.ietf.org/drafts/current/>.

Internet-Drafts are draft documents valid for a maximum of six months and may be updated, replaced, or obsoleted by other documents at any time. It is inappropriate to use Internet-Drafts as reference material or to cite them other than as "work in progress."

This Internet-Draft will expire on August 18, 2014.

Copyright Notice

Copyright (c) 2014 IETF Trust and the persons identified as the document authors. All rights reserved.

This document is subject to BCP 78 and the IETF Trust's Legal Provisions Relating to IETF Documents (<http://trustee.ietf.org/license-info>) in effect on the date of

publication of this document. Please review these documents carefully, as they describe your rights and restrictions with respect to this document. Code Components extracted from this document must include Simplified BSD License text as described in Section 4.e of the Trust Legal Provisions and are provided without warranty as described in the Simplified BSD License.

Table of Contents

1. Introduction	2
2. Terminology	3
3. Basic Structure of Test cases	3
4. Basic Test Cases	7
4.1. Variable Available Capacity	7
4.2. Maximum Media Bit Rate is Greater than Link Capacity	10
4.3. Competing Flows with same RMCAT Algorithm	13
4.4. RMCAT Flow competing with a long TCP Flow	16
4.5. RMCAT Flow competing with short TCP Flows	19
4.6. Congested Feedback Link	22
4.7. Round Trip Time Fairness	27
4.8. Media Pause and Resume	29
4.9. Explicit Congestion Notification Usage	31
5. Wireless Access Links	31
5.1. Cellular Network Specific Test Cases	31
5.2. Wi-Fi Network Specific Test Cases	31
6. Security Considerations	31
7. IANA Considerations	31
8. Acknowledgements	31
9. References	32
9.1. Normative References	32
9.2. Informative References	33
Authors' Addresses	33

1. Introduction

This memo describes a set of test cases for evaluating candidate RMCAT congestion control algorithm proposals, it is based on the guidelines enumerated in [I-D.ietf-rmcat-eval-criteria] and requirements discussed in [I-D.ietf-rmcat-cc-requirements]. The test cases cover basic usage scenarios and are described using a common structure, which allows for additional test cases to be added to those described herein to accommodate other topologies and/or the modeling of different link characteristics. Each test case incorporates the metrics, evaluation guidelines and parameters described in [I-D.ietf-rmcat-eval-criteria]. It is the intention of this work to capture the consensus of the RMCAT working group participants regarding the test cases upon which the performance of the candidate RMCAT proposals should be evaluated.

2. Terminology

The terminology defined in RTP [RFC3550], RTP Profile for Audio and Video Conferences with Minimal Control [RFC3551], RTCP Extended Report (XR) [RFC3611], Extended RTP Profile for RTCP-based Feedback (RTP/AVPF) [RFC4585], Support for Reduced-Size RTCP [RFC5506], and RTP Circuit Breaker algorithm [I-D.ietf-avtcore-rtp-circuit-breakers] apply.

3. Basic Structure of Test cases

All test cases in this document follow a basic structure allowing implementers to describe a new test scenarios without repeatedly explaining common attributes. The structure includes a general description section that describes the test case and motivations, additionally it defines a set of attributes that characterize the testbed, i.e., the network path between communicating peers and the diverse traffic sources.

o Define the test case:

- * General description: describes the motivation and the goals of the test case.
- * Additionally, describe the desired rate adaptation behaviour.
- * Define a checklist to evaluate the desired behaviour: this indicates the minimum set of metrics (e.g., link utilization, media sending rate) that a proposed algorithm needs to measure to validate the expected rate adaptation behaviour. It should also indicate the time granularity (e.g., averaged over 10ms, 100ms, or 1s) for measuring certain metrics. Typical measurement interval is 200ms.

- o Define testbed topology: every test case needs to define an evaluation testbed topology. Figure 1 shows such an evaluation topology. In this evaluation topology, S1..Sn are traffic sources, these sources either generate media traffic and use an RMCAT candidate congestion control algorithm or use a different congestion control algorithm than that one used in the media source. R1..Rn are the corresponding receivers. A test case can have one or more such traffic source (S) and corresponding receiver (R). The path from the source to destination is denoted as upstream and the path from a destination to a source is denoted as downstream. It is possible for a testbed to have different path characteristics in either directions. The following basic structure of test case has been described from the perspective of endpoints attached on the left-hand side of Figure 1.

- + Bottleneck queue type: for example, Droptail, FQ-CoDel, or PIE.
- + Bottleneck queue size: defines size of queue in terms of queuing time when the queue is full (in milliseconds).
- + Path loss ratio: characterize the non-congested losses observed on the end-to-end path. MUST describe the loss pattern or loss model.
- * Application-related: defines the media-related behaviour for implementing the test case
 - + Media Source: defines the characteristics of the media sources present. When using more than one media source, the different attributes are enumerated separately for each different media source.
 - Media type: Video/Voice/Application/Text
 - Media flow direction: upstream, downstream or both.
 - Number of media sources: defines the total number of media sources
 - Media codec: Constant Bit Rate (CBR) or Variable Bit Rate (VBR)
 - Media source behaviour: describes the media encoder behavior. It MUST define the main parameters that affect the adaptation behaviour. This may include but not limited to:
 - o Adaptability: describes the adaptation options. For example, in case of video it defines the range of bitrate adaptation, frame rate adaptation, video resolution adaptation. In case of voice, it defines the range of bitrate adaptation, sampling rate, frame size.
 - o Output variation : for VBR encoder, it defines the encoder output variation from the average target rate. For example on average the encoder output may vary between 5% to 15% above or below the average target bit rate.
 - o Encoder's responsiveness to a new bit rate request: value typically between 10ms to 1000ms.

- Media content: describes the chosen media sequences; For example, test sequences are available at: [xiph-seq] [HEVC-seq].
 - Media timeline: describes the point when the media source is introduced and removed from the testbed. For example, the media source may begin transmitting when the test case begins or a few seconds after, etc.
 - Startup behaviour: the media starts at a defined bit rate, which may be the minimum, maximum bit rate, or a value in between (in Kbps).
- + Competing traffic source: describes the characteristics of the competing traffic source, the different types of competing flows are enumerated in [I-D.ietf-rmcat-eval-criteria].
- Traffic direction: Upstream, downstream or both.
 - Type of sources: defines the types of competing traffic sources. Types of competing traffic flows are listed in [I-D.ietf-rmcat-eval-criteria]. For example, the number of TCP flows connected to a web browser, the mean size and distribution of the content downloaded.
 - Number of sources: defines the total number of competing sources of each type.
 - Congestion control: enumerate the congestion control used by each type of competing traffic.
 - Traffic timeline: describes when the competing traffic is added and removed from the test case.
- * Additional attributes: describes attributes essential for implementing a test case which are not included in the above structure. These attributes MUST be well defined, so that other implementers are able to implement it.

Any attribute can have a set of values (enclosed within "[]"). Each member value of such a set MUST be treated as different value for the same attribute. It is desired to run separate tests for each such attribute value.

The test cases described in this document follow the above structure.

4. Basic Test Cases

4.1. Variable Available Capacity

In this test case the bottleneck-link capacity between the two endpoints varies over time. This test is designed to measure the responsiveness of the candidate algorithm. This test tries to address the requirement 1(a) in [I-D.ietf-rmcat-cc-requirements], which requires the algorithm to adapt the flow(s) and provide lower end-to-end latency when there exists:

- o an intermediate bottleneck
- o change in available capacity due to interface change and/or routing change.
- o persistent network load due to competing traffic

It should be noted that the exact variation in available capacity due to any of the above depends on the under-lying technologies. Hence, we describe a set of known factors, which may be extended to devise a more specific test case targeting certain behaviour in a certain network environment.

Expected behavior: the candidate algorithms is expected to detect the variation in available capacity and adapt the media stream(s) accordingly. The flows stabilize around their maximum bitrate as the as the maximum link capacity is large enough to accommodate the flows. When the available capacity drops, the flow(s) adapts by decreasing its sending bit rate, and when congestion disappears, the flow(s) are again expected to ramp up.

To evaluate the performance of the candidate algorithms it is expected to log enough information to visualize the following metrics at a fine enough time granularity:

1. Flow level:
 - a. End-to-end delay
 - b. Losses observed at the receiving endpoint
 - c. Feedback message overhead
2. Transport level:
 - a. Bandwidth utilization

b. Queue length (ms):

- + average over the length of the session
- + 5 and 95 percentile
- + median, maximum, minimum

Testbed Topology: Two (2) media sources S1 and S2 are connected to their corresponding R1 and R2. The media traffic is transported over the upstream path and corresponding feedback/control traffic is transported over the downstream path.

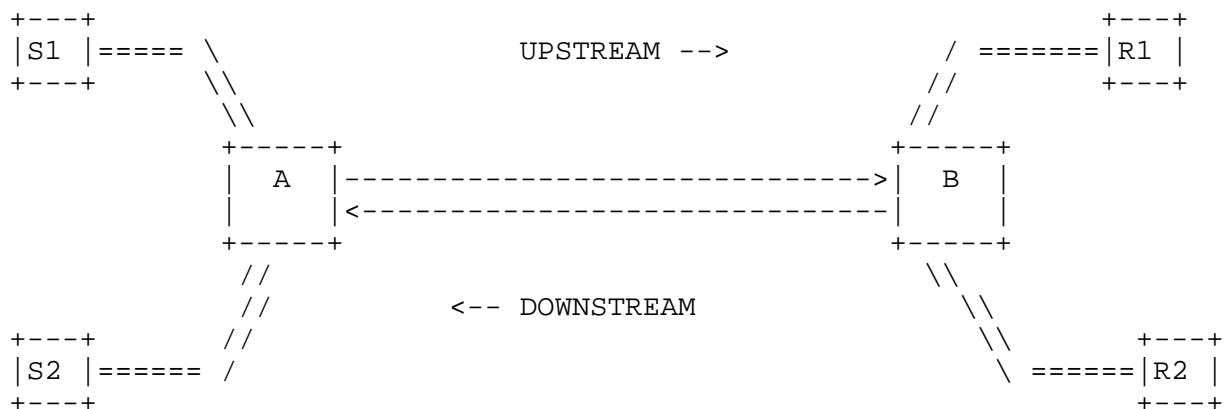


Figure 2: Testbed Topology for Variable Available Capacity

Testbed attributes:

- o Test duration: 120s
- o Path characteristics:
 - * Path direction: Upstream and downstream.
 - * Bottleneck-link capacity: 4Mbps.
 - * One-Way propagation delay: [10ms, 100ms].
 - * Maximum end-to-end jitter: 30ms.
 - * Bottleneck queue type: Drop tail.
 - * Bottleneck queue size: 300ms.
 - * path loss ratio: 0%.

- o Application-related:
 - * Media Source:
 - + Media type: Video
 - + Media direction: Upstream.
 - + Number of media sources: Two (2).
 - + Media codec: VBR
 - + Media source behaviour:
 - Adaptability:
 - o Bit rate range: 150 Kbps - 1.5 Mbps
 - o Frame Resolution: 144p - 720p (or 1080p)
 - o Frame rate: 10fps - 30fps
 - Variation from target bitrate: +/-5%
 - Responsiveness to new bit rate request: 100ms
 - + Media content: Foreman media sequence.
 - + Media timeline:
 - Start time: 0s.
 - End time: 119s.
 - * Media startup behaviour: [200Kbps, 1500Kbps].
 - * Competing traffic:
 - + Number of sources : Zero (0)
- o Test specific setup:
 - * Number of bandwidth variations: Three (4)
 - * Variation pattern:
 - + Sequence number: 1

- + Path direction: Upstream
- + Bottleneck Capacity: 2Mbps.
- + Start time: 30s
- * Variation pattern:
 - + Sequence number: 2
 - + Path direction: Upstream
 - + Bottleneck Capacity: 3.5Mbps
 - + Start time: 50s
- * Variation pattern:
 - + Sequence number: 3
 - + Path direction: Upstream
 - + Bottleneck Capacity: 1Mbps
 - + Start time: 70s
- * Variation pattern:
 - + Sequence number: 4
 - + Path direction: Upstream
 - + Bottleneck Capacity: 2Mbps
 - + Start time: 100s

4.2. Maximum Media Bit Rate is Greater than Link Capacity

In this case, the application will attempt to ramp up to its maximum bit rate, since the link capacity is limited to a value lower, the congestion control scheme is expected to stabilize the sending bit rate close to the available bottleneck capacity. This situation can occur when the endpoints are connected via thin long networks even though the advertised capacity of the access network may be higher. The test case addresses the requirement 1 and 10 of the [I-D.ietf-rmcat-cc-requirements].

Expected behavior: the candidate algorithm is expected to detect the path capacity constraint, converges to bottleneck link's capacity and adapt the flow to avoid unwanted oscillation when the sending bit rate is approaching the bottleneck link's capacity. The oscillations occur when the media flow(s) attempts to reach its maximum bit rate, overshoots the usage of the available bottleneck capacity, to rectify it reduces the bit rate and starts to ramp up again.

Testbed topology: One media source S1 is connected to corresponding R1. The media traffic is transported over the upstream path and corresponding feedback/control traffic is transported over the downstream path.




Figure 3: Testbed Topology for Limited Link Capacity

To evaluate the performance of the candidate algorithms it is expected to log enough information to visualize the following metrics at a fine enough time granularity:

1. Flow level:
 - a. End-to-end delay.
 - b. RTP packet losses observed at the receiving endpoint.
 - c. Variation in sending bit rate and goodput. Mainly observing the frequency and magnitude of oscillations.
 - d. Convergence time.
 - e. Feedback message overhead.
2. Transport level:
 - a. Bandwidth utilization.
 - b. Queue length (ms):
 - + average over the length of the session
 - + 5 and 95 percentile

+ median

Testbed attributes:

- o Test duration: 60s
- o Path characteristics:
 - * Path direction: Upstream and downstream.
 - * Bottleneck-link capacity: 1Mbps
 - * One-Way propagation delay: [10ms,100 ms]
 - * Maximum end-to-end jitter: 30ms.
 - * Bottleneck queue type: Droptail. Additional tests with other AQM schemes are recommended: FQ-CoDel, PIE
 - * Bottleneck size: 300ms
 - * Path loss ratio: 0%
- o Application-related:
 - * Media Source:
 - + Media type: Video
 - + Media direction: Upstream.
 - + Number of media sources: One (1)
 - + Media codec: VBR
 - + Media source behaviour:
 - Adaptability:
 - o Bit rate range: 150 Kbps - 1.5 Mbps
 - o Frame Resolution: 144p - 720p (or 1080p)
 - o Frame rate: 10fps - 30fps
 - Variation from target bitrate: +/-5%
 - Responsiveness to new bit rate request: 100ms

- + Media content: Foreman video sequence
- + Media timeline:
 - Start time: 0s.
 - End time: 59s.
- * Media startup behaviour: [200Kbps, 1500Kbps].
- * Competing traffic:
 - + Number of sources : Zero (0)

o Test specific setup: None

4.3. Competing Flows with same RMCAT Algorithm

In this test case, more than one RMCAT media flow shares the bottleneck link and each of them uses the same congestion control algorithm. This is a typical scenario wherein a real-time interactive application sends more than one media flows to the same destination and these flows are multiplexed over the same port. In such a scenario it is likely that the flows will be routed via the same path and need to share the available bandwidth amongst themselves. For the sake of simplicity it is assumed that there are no other competing traffic sources in the bottleneck link and that there is sufficient capacity to accommodate all the flows individually. While this appears to be a variant of the test case defined in Section 4.1 , however it tests the capacity sharing distribution of the candidate algorithm. Whereas, the previous test case measures the stability and responsiveness of the candidate algorithm. This test case particularly addresses the requirements 2,3 and 10 in [I-D.ietf-rmcat-cc-requirements].

Expected behavior: It is expected that the competing flows will converge to an optimum bit rate to accommodate all the flows with minimum possible latency and loss. Specifically, the test introduces three media flows at different time instances, when the second flow appears there should still be room to accommodate another flow on the bottleneck link. Lastly, when the third flow appears the bottleneck link should be saturated.

To evaluate the performance of the candidate algorithms it is expected to log enough information to visualize the following metrics metrics at a fine enough time granularity:

1. Flow level:

- a. End-to-end delay.
 - b. RTP packet losses observed at the receiving endpoint.
 - c. Variation in sending bit rate and goodput. Mainly observing the frequency and magnitude of oscillations.
 - d. Convergence time.
 - e. Feedback message overhead.
2. Transport level:
- a. Bandwidth utilization.
 - b. Queue length (ms):
 - + average over the length of the session
 - + 5 and 95 percentile
 - + median

Testbed topology: Three media sources S1, S2, S3 are connected to respective R1, R2, R3. The media traffic is transported over the upstream path and corresponding feedback/control traffic is transported over the downstream path.




Figure 4: Testbed Topology for Multiple RMCAT Flows

Testbed attributes:

- o Test duration: 60s

- o Path characteristics:
 - * Path direction: Upstream, Downstream
 - * Bottleneck-link capacity: 3.5Mbps
 - * One-Way propagation delay: [10ms, 50ms]
 - * Maximum end to end jitter: 30ms
 - * Bottleneck queue type: Droptail
 - * Bottleneck queue size: 300ms
 - * Link loss ratio: 0.0%
- o Application-related:
 - * Media Source:
 - + Media type: Video
 - + Media direction: Upstream
 - + Number of media sources: Three (3)
 - + Media codec: VBR
 - + Media source behaviour:
 - Adaptability:
 - o Bit rate range: 150 Kbps - 1.5 Mbps
 - o Frame Resolution: 144p - 720p (or 1080p)
 - o Frame rate: 10fps - 30fps
 - Variation from target bitrate: +/-5%
 - Responsiveness to new bit rate request: 100ms
 - + Media content: Foreman video sequence
 - + Media timeline: New media flows are added sequentially, at short time intervals. See test specific setup below.
 - + Media startup behaviour: 200Kbps.

- * Competing traffic:
 - + Number of sources : Zero (0)
- o Test specific setup:
 - * Media flow timeline:
 - + Flow no: One (1)
 - + Start time: 0s
 - + End time: 59s
 - * Media flow appearance:
 - + Flow no: Two (2)
 - + Start time: 20s
 - + End time: 59s
 - * Media flow appearance:
 - + Flow no: Three (3)
 - + Start time: 40s
 - + End time: 59s

4.4. RMCAT Flow competing with a long TCP Flow


In this test case, one or more RMCAT media flow shares the bottleneck link with at least one long lived TCP flows. Long lived TCP flows download data throughout the session and are expected to have infinite amount of data to send and receive. This is a scenario wherein a multimedia application co-exists with a large file download. The test case measures the adaptivity of the candidate algorithm to competing traffic, it addresses the requirement 8 in [I-D.ietf-rmcat-cc-requirements].

Expected behavior: depending on the convergence observed in test case 4.1 and 4.2, the candidate algorithm may be able to avoid congestion collapse. In the worst case, the media stream will fall to the minimum media bit rate.

To evaluate the performance of the candidate algorithms it is expected to log enough information to visualize the following metrics at a fine enough time granularity:

1. Flow level:
 - a. End-to-end delay for the RMCAT flow.
 - b. RTP packet losses observed at the receiving endpoint.
 - c. Variation in sending bit rate and goodput. Mainly observing the frequency and magnitude of oscillations.
 - d. Variation in the sending rate of the TCP flow
 - e. TCP throughput.
 - f. Convergence time.
 - g. Feedback message overhead.
2. Transport level:
 - a. Bandwidth utilization.
 - b. Queue length (ms):
 - + average over the length of the session
 - + 5 and 95 percentile

Testbed topology: One (1) media source S1 is connected to corresponding R1, but both endpoints are additionally receiving and sending data, respectively. The media traffic (S1->S2) is transported over the upstream path and corresponding feedback/control traffic is transported over the downstream path. Likewise media traffic (S2->S1) is transported over the downstream path and corresponding feedback/control traffic is transported over the upstream path. The TCP uploads traffic over upstream path and downloads over downstream path. Hence, TCP traffic exits in both path directions.






Figure 5: Testbed Topology for TCP vs RMCAT Flows

Testbed attributes:

- o Test duration: 120s
- o Path characteristics:
 - * Path direction: Upstream, Downstream
 - * Bottleneck-link capacity: 2Mbps
 - * One-Way propagation delay: [10ms, 150ms]
 - * Maximum end to end jitter: 30ms
 - * Bottleneck queue type: Droptail, but would benefit from running the same test with different AQM schemes: FQ-Codel, or PIE.
 - * Bottleneck queue size: [20ms, 250ms, 1000ms]
 - * Path loss ratio: 0.0%
- o Application-related:
 - * Media Source:
 - + Media type: Video
 - + Media direction: Upstream and Downstream
 - + Number of media sources: One (1)
 - + Media codec: VBR
 - + Media source behaviour:
 - Adaptability:
 - o Bit rate range: 150 Kbps - 1.5 Mbps

- o Frame Resolution: 144p - 720p (or 1080p)
 - o Frame rate: 10fps - 30fps
 - Variation from target bitrate: +/-5%
 - Responsiveness to new bit rate request: 100ms
 - + Media content: Foreman video sequence
 - + Media timeline:
 - Start time: 5s.
 - End time: 119s.
 - + Media startup behaviour: [200Kbps, 1500Kbps].
- * Competing traffic:
 - + Number and Types of sources : one (1), long-lived TCP
 - + Traffic direction : Downstream
 - + Congestion control: Default TCP congestion control.
 - + Traffic timeline:
 - Start time: 0s.
 - End time: 119s.
- o Test specific setup: None

4.5. RMCAT Flow competing with short TCP Flows

In this test case, one or more RMCAT media flow shares the bottleneck link with at multiple short-lived TCP flows. Short-lived TCP flows resemble the on/off pattern observed in the web traffic, wherein clients (browsers) connect to a server and download a resource (typically a webpage, few images, text files, etc.) using several TCP connections (up to 4). This scenario shows the performance of the multimedia application when several browser windows are active. The test case measures the adaptivity of the candidate algorithm to competing web traffic, it addresses the requirements 2 in [I-D.ietf-rmcat-cc-requirements].

Depending on the number of short TCP flows, the cross-traffic either appears as a short burst flow or resembles a long TCP flow. The intention of this test is to observe the impact of short-term burst on the behaviour of the candidate algorithm.

To evaluate the performance of the candidate algorithms it is expected to log enough information to visualize the following metrics at a fine enough time granularity:

1. Flow level:
 - a. End-to-end delay for the RMCAT flow.
 - b. RTP packet losses observed at the receiving endpoint.
 - c. Variation in sending bit rate and goodput. Mainly observing the frequency and magnitude of oscillations.
 - d. Variation in the sending rate of the TCP flow.
 - e. TCP throughput.
 - f. Convergence time.
 - g. Feedback message overhead.
2. Transport level:
 - a. Bandwidth utilization.
 - b. Queue length (ms):
 - + average over the length of the session
 - + 5 and 95 percentile

Testbed topology: One (1) media source S1 is connected to corresponding R1, but both endpoints are additionally receiving and sending data, respectively. The media traffic (S1->S2) is transported over the upstream path and corresponding feedback/control traffic is transported over the downstream path. Likewise media traffic (S2->S1) is transported over the downstream path and corresponding feedback/control traffic is transported over the upstream path. The TCP uploads traffic over upstream path and downloads over downstream path. Hence, TCP traffic exits in both path directions. The topology described here is similar to the one described in Figure 5.

Testbed attributes:

- o Test duration: 300s
- o Path characteristics:
 - * Path direction: Upstream, Downstream
 - * Bottleneck-link capacity: 2.0Mbps
 - * One-Way propagation delay: [10ms, 150ms]
 - * Maximum end to end jitter: 30ms
 - * Bottleneck queue type: Droptail, but would benefit from running the same test with different AQM schemes: FQ-Codel, or PIE.
 - * Bottleneck queue size: 300ms
 - * Path loss ratio: 0.0%
- o Application-related:
 - * Media Source:
 - + Media type: Video
 - + Media direction: Upstream and Downstream
 - + Number of media sources: One (1)
 - + Media codec: VBR
 - + Media source behaviour:
 - Adaptability:
 - o Bit rate range: 150 Kbps - 1.5 Mbps
 - o Frame Resolution: 144p - 720p (or 1080p)
 - o Frame rate: 10fps - 30fps
 - Variation from target bitrate: +/-5%
 - Responsiveness to new bit rate request: 100ms
 - + Media content: Foreman video sequence

- + Media timeline:
 - Start time: 0s.
 - End time: 299s.
- + Media startup behaviour: [200Kbps, 1500Kbps].
- * Competing traffic:
 - + Number and Types of sources : Ten (10), short-lived TCP flows.
 - + Traffic direction : Downstream
 - + Congestion algorithm: Default TCP Congestion control.
 - + Traffic timeline: Each short TCP flow is modeled as a sequence of file downloads interleaved with idle periods. See test specific setup. Not all short TCPs start at the same time, 2 start in the ON state while 8 start in an OFF state. The model for the idle times for the OFF state is discussed in the Short-TCP model.
- o Test specific setup:
 - * Short-TCP traffic model:
 - + File sizes: uniform distribution between 100KB to 1MB
 - + Idle period: the duration of the OFF state is derived from an exponential distribution with the mean value of 10 seconds.

4.6. Congested Feedback Link

RMCAT WG has been chartered to define algorithms for RTP hence it is assumed that RTCP, RTP header extension or such would be used as signalling means for the adaptation algorithm in the backchannel. Due to asymmetry nature of the link between communicating peers it is possible to observe lack such backchannel information due to impaired backchannel link (even when forward channel might be unimpaired). This test case is designed to observe candidate congestion control behaviour in such an event. This test case addresses requirement number 5 and in particular, requirement number 7.

It is expected that the candidate algorithms should cope with the lack of backchannel information and adapt to minimize the performance degradation of media flows in the forward channel.


It SHOULD be noted that for this test case log is needed for the reference case where the downstream channel has no impairments

To evaluate the performance of the candidate algorithms it is expected to log enough information to visualize the following metrics at a fine enough time granularity:

1. Flow level:
 - a. End-to-end delay
 - b. Losses observed at the receiving endpoint
 - c. Feedback message overhead
2. Transport level:
 - a. Bandwidth utilization
 - b. Queue length (ms):
 - + average over the length of the session
 - + 5 and 95 percentile

•

Testbed topology: One (1) media source S1 is connected to corresponding R1, but both endpoints are additionally receiving and sending data, respectively. The media traffic (S1->S2) is transported over the upstream path and corresponding feedback/control traffic is transported over the downstream path. Likewise media traffic (S2->S1) is transported over the downstream path and corresponding feedback/control traffic is transported over the upstream path.






Figure 6: Testbed Topology for Congested Feedback Link

Testbed attributes:

- o Test duration: 60s
- o Path characteristics:
 - * Path direction: Upstream and downstream.
 - * Bottleneck-link capacity: 2Mbps.
 - * One-Way propagation delay: [10ms, 100ms].
 - * Maximum end-to-end jitter: 30ms.
 - * Bottleneck queue type: Drop tail.
 - * Bottleneck queue size: 300ms.
 - * path loss ratio: 0%.
- o Application-related:
 - * Media Source:
 - + Media type: Video
 - + Media direction: Upstream
 - + Number of media sources: One (1).
 - + Media codec: VBR
 - + Media source behaviour:
 - Adaptability:
 - o Bit rate range: 150 Kbps - 1.5 Mbps
 - o Frame Resolution: 144p - 720p (or 1080p)

- o Frame rate: 10fps - 30fps
 - Variation from target bitrate: +/-5%
 - Responsiveness to new bit rate request: 100ms
- + Media content: Foreman media sequence
- + Media timeline:
 - Start time: 0s.
 - End time: 59s.
- * Media Source:
 - + Media type: Video
 - + Media direction: Downstream
 - + Number of media sources: One (1).
 - + Media codec: VBR
 - + Media source behaviour:
 - Adaptability:
 - o Bit rate range: 150 Kbps - 1.5 Mbps
 - o Frame Resolution: 144p - 720p (or 1080p)
 - o Frame rate: 10fps - 30fps
 - Variation from target bitrate: +/-5%
 - Responsiveness to new bit rate request: 100ms
 - + Media content: Foreman media sequence
 - + Media timeline:
 - Start time: 0s.
 - End time: 59s.
- * Media startup behaviour: 200Kbps.

- * Competing traffic:
 - + Number of sources : Zero (0)
- o Test specific setup:
 - * Number of bandwidth variation: Two (2)
 - * Variation pattern:
 - + Sequence number: 1
 - + Path direction: Upstream
 - + Bottleneck Capacity: 1Mbps
 - + Start time: 20s
 - * Variation pattern:
 - + Sequence number: 2
 - + Path direction: Upstream
 - + Bottleneck Capacity: 600Kbps
 - + Start time: 30s
 - * Variation pattern:
 - + Sequence number: 3
 - + Path direction: Upstream
 - + Bottleneck Capacity: 2Mbps
 - + Start time: 50s
 - * Variation pattern:
 - + Sequence number: 1
 - + Path direction: Downstream
 - + Bottleneck Capacity: 800kbps
 - + Start time: 30s

- * Variation pattern:
 - + Sequence number: 2
 - + Path direction: Downstream
 - + Bottleneck Capacity: 2Mbps
 - + Start time: 50s

4.7. Round Trip Time Fairness

In this test case, more than one RMCAT media flow shares the bottleneck link, but the end-to-end path latency for each RMCAT flow is different. For the sake of simplicity it is assumed that there are no other competing traffic sources in the bottleneck link and that there is sufficient capacity to accommodate all the flows. While this appears to be a variant of the test case 4.2, it tests the capacity sharing distribution of the candidate algorithm under different RTTs. This test case particularly addresses the requirements 2 [I-D.ietf-rmcat-cc-requirements].

It is expected that the competing flows will converge to an optimum bit rate to accommodate all the flows with minimum possible latency and loss. Specifically, the test introduces five media flows at the same time instance.

To evaluate the performance of the candidate algorithms it is expected to log enough information to visualize the following metrics at a fine enough time granularity:

1. Flow level:
 - a. End-to-end delay.
 - b. RTP packet losses observed at the receiving endpoint.
 - c. Variation in sending bit rate and goodput. Mainly observing the frequency and magnitude of oscillations.
 - d. Convergence time.
2. Transport level:
 - a. Bandwidth utilization.
 - b. Queue length (ms):

- + average over the length of the session
- + 5 an 95 percentile

Testbed Topology: Five (5) media sources S1,S2,...,S5 are connected to their corresponding media sinks R1,R2,...,R5. The media traffic is transported over the upstream path and corresponding feedback/control traffic is transported over the downstream path. The topology is the same as test case 4.3 defined in Section 4.3. The end-to-end path delay for S1-R1 is 10ms, S2-R2 is 25ms, S3-R3 is 50ms, S4-R4 is 100ms, S5-R5 is 150ms.

Testbed attributes:

- o Test duration: 60s
- o Path characteristics:
 - * Path direction: Upstream, Downstream
 - * Number of bottlenecks: One (1)
 - * Bottleneck link capacity: 3.0Mbps
 - * One-Way propagation delay for each path is: 10ms, 25ms, 50ms, 100ms, 150ms.
 - * Maximum end to end jitter: 30ms
 - * Bottleneck queue type: Droptail
 - * Bottleneck queue size: 300ms
 - * Link loss ratio: 0.0%
- o Application-related:
 - * Media Source:
 - + Media direction: Upstream
 - + Number of media sources: Five (5)
 - + Encoder configuration:
 - Bit rate generation: VBR
 - Bit rate range: 150 Kbps - 1.5 Mbps

- Frame Resolution: 144p-720p (or 1080p)
- Frame rate: 10fps-30fps
- Variation from target bit rate: +/-5%
- Responsiveness to new bit rate request: 60ms
- + Media content: Foreman video sequence
- + Media timeline:
 - Start time: 0s.
 - End time: 59s.
- + Media startup behaviour: [200 Kbps, 1500 Kbps].
- * Competing traffic:
 - + Number of sources : Zero (0)
- o Test specific setup: None

4.8. Media Pause and Resume

In this test case, more than one real-time interactive media flows share the link bandwidth and all flows reach to a steady state by utilizing the link capacity in an optimum way. At these stage one of the media flow is paused for a moment. This event will result in more available bandwidth for the rest of the flows and as they are on a shared link. When the paused media flow will resume it would no longer have the same bandwidth share on the link. It has to make it way through the other existing flows in the link to achieve a fair share of the link capacity. This test case is important specially for real-time interactive media which consists of more than one media flows and can pause/resume media flow at any point of time during the session. This test case directly addresses the requirement number 1.B in [I-D.ietf-rmcat-cc-requirements]. One can think it as a variation of test case 4.3. However, it is different as the candidate algorithms can use different strategies to increase it s efficiency, for example the fairness, convergence time, reduce oscillation etc, by capitalizing the fact that they have previous information of the link.

To evaluate the performance of the candidate algorithms it is expected to log enough information to visualize the following metrics at a fine enough time granularity:

1. Flow level:
 - a. End-to-end delay.
 - b. RTP packet losses observed at the receiving endpoint.
 - c. Variation in sending bit rate and goodput. Mainly observing the frequency and magnitude of oscillations.
 - d. Convergence time.
 - e. Feedback message overhead.
2. Transport level:
 - a. Bandwidth utilization.
 - b. Queue length (ms):
 - + average over the length of the session
 - + 5 and 95 percentile

Testbed Topology: Same as test case 4.3 defined in Section 4.3

Testbed attributes: The general description of the testbed parameters are same as test case 4.3 with changes in the test specific setup as below-

- o Other test specific setup:
 - * Media flow timeline:
 - + Flow no: One (1)
 - + Start time: 0s
 - + Flow duration: 59s
 - + Pause time: not required
 - + Resume time: not required
 - * Media flow appearance:
 - + Flow no: Two (2)
 - + Start time: 0s

- + Flow duration: 59s
- + Pause time: 20s
- + Resume time: 30s
- * Media flow appearance:
 - + Flow no: Three (3)
 - + Start time: 0s
 - + Flow duration: 59s
 - + Pause time: not required
 - + Resume time: not required

4.9. Explicit Congestion Notification Usage

TBD

5. Wireless Access Links

5.1. Cellular Network Specific Test Cases

Additional cellular network specific test cases are define in
[I-D.draft-sarker-rmcat-cellular-eval-test-cases]

5.2. Wi-Fi Network Specific Test Cases

TBD

[Editor's Note: We should encourage people to come up with possible
WiFi Network specific test cases]

6. Security Considerations

Security issues have not been discussed in this memo.

7. IANA Considerations

There are no IANA impacts in this memo.

8. Acknowledgements

Much of this document is derived from previous work on congestion
control at the IETF.

The content and concepts within this document are a product of the discussion carried out in the Design Team.

9. References

9.1. Normative References

- [RFC3550] Schulzrinne, H., Casner, S., Frederick, R., and V. Jacobson, "RTP: A Transport Protocol for Real-Time Applications", STD 64, RFC 3550, July 2003.
- [RFC3551] Schulzrinne, H. and S. Casner, "RTP Profile for Audio and Video Conferences with Minimal Control", STD 65, RFC 3551, July 2003.
- [RFC3611] Friedman, T., Caceres, R., and A. Clark, "RTP Control Protocol Extended Reports (RTCP XR)", RFC 3611, November 2003.
- [RFC4585] Ott, J., Wenger, S., Sato, N., Burmeister, C., and J. Rey, "Extended RTP Profile for Real-time Transport Control Protocol (RTCP)-Based Feedback (RTP/AVPF)", RFC 4585, July 2006.
- [RFC5506] Johansson, I. and M. Westerlund, "Support for Reduced-Size Real-Time Transport Control Protocol (RTCP): Opportunities and Consequences", RFC 5506, April 2009.
- [RFC5506] Johansson, I. and M. Westerlund, "Support for Reduced-Size Real-Time Transport Control Protocol (RTCP): Opportunities and Consequences", RFC 5506, April 2009.
- [I-D.ietf-avtcore-rtp-circuit-breakers]
Perkins, C. and V. Singh, "RTP Congestion Control: Circuit Breakers for Unicast Sessions", draft-ietf-avtcore-rtp-circuit-breakers-01 (work in progress), October 2012.
- [I-D.ietf-rmcat-eval-criteria]
Singh, V. and J. Ott, "Evaluating Congestion Control for Interactive Real-time Media", draft-ietf-rmcat-eval-criteria-00 (work in progress), January 2014.
- [I-D.ietf-rmcat-cc-requirements]
Jesup, R., "Congestion Control Requirements For RMCAT", draft-ietf-rmcat-cc-requirements-00 (work in progress), July 2013.
- [I-D.draft-sarker-rmcat-cellular-eval-test-cases]

Sarker, Z., "Evaluation Test Cases for Interactive Real-Time Media over Cellular Networks", , <<http://www.ietf.org/id/draft-sarker-rmcat-cellular-eval-test-cases-00.txt>>.

9.2. Informative References

- [I-D.ietf-rtcweb-use-cases-and-requirements]
Holmberg, C., Hakansson, S., and G. Eriksson, "Web Real-Time Communication Use-cases and Requirements", draft-ietf-rtcweb-use-cases-and-requirements-10 (work in progress), December 2012.
- [RFC5033] Floyd, S. and M. Allman, "Specifying New Congestion Control Algorithms", BCP 133, RFC 5033, August 2007.
- [RFC5166] Floyd, S., "Metrics for the Evaluation of Congestion Control Mechanisms", RFC 5166, March 2008.
- [RFC5681] Allman, M., Paxson, V., and E. Blanton, "TCP Congestion Control", RFC 5681, September 2009.
- [SA4-EVAL]
R1-081955, 3GPP., "LTE Link Level Throughput Data for SA4 Evaluation Framework", 3GPP R1-081955, 5 2008.
- [SA4-LR] S4-050560, 3GPP., "Error Patterns for MBMS Streaming over UTRAN and GERAN", 3GPP S4-050560, 5 2008.
- [xiph-seq]
Xiph.org, , "Video Test Media",
<http://media.xiph.org/video/derf/> , .
- [HEVC-seq]
HEVC, , "Test Sequences",
http://www.netlab.tkk.fi/~varun/test_sequences/ , .
- [TCP-eval-suite]
Lachlan, A., Marcondes, C., Floyd, S., Dunn, L., Guillier, R., Gang, W., Eggert, L., Ha, S., and I. Rhee, "Towards a Common TCP Evaluation Suite", Proc. PFLDnet. 2008, August 2008.

Authors' Addresses

Zaheduzzaman Sarker
Ericsson AB
Luleae, SE 977 53
Sweden

Phone: +46 10 717 37 43
Email: zaheduzzaman.sarker@ericsson.com

Varun Singh
Aalto University
School of Electrical Engineering
Otakaari 5 A
Espoo, FIN 02150
Finland

Email: varun@comnet.tkk.fi
URI: <http://www.netlab.tkk.fi/~varun/>

Xiaoqing Zhu
Cisco Systems
510 McCarthy Blvd
Milpitas, CA 95134
USA

Email: xiaoqzhu@cisco.com

Michael A. Ramalho
Cisco Systems, Inc.
8000 Hawkins Road
Sarasota, FL 34241
USA

Phone: +1 919 476 2038
Email: mramalho@cisco.com